


MATHGUARD

A control layer for AI

Policy before content or effects.



One shared enforcement path



Prompts Model replies MCP Agent messages Artifacts

AI cannot relax a hard denial

One request. Two independent checks.

RULE	LOCAL AI	FINAL RESULT
Signature matched	“Safe”	Block
No hard denial	Threat detected	Block / review

Pass, redact, review or block.

Live policy, safe reload

EDIT

GATEWAY RESPONSE

Valid policy

Activate

Invalid edit

Keep last valid

No valid policy

Block all traffic

No restart. No quota reset.

Checks before content escapes

CONTENT

CONTROL

Secrets

Block or redact

Personal data

Redact or block

Executable / unpinned artifact

Block

Bounded parsing. Hash pins. No deserialization.

Budget before dispatch

Reserve the upper bound

Tokens Time Calls

Exhausted budget stops the next call.

Timeout keeps the charge and triggers quarantine.

Approval and safe retry



Restart restores acknowledged state.

An exact retry adds no second debit.

A clear operator view

OPERATOR NEED

VISIBLE EVIDENCE

What stopped?

Decision + reason

What remains?

Resource headroom

What changed?

Policy / feed / posture

Sanitized logs and management export.

Evidence with a defined scope

156

AXIOM AUDIT RECORDS

Compiled Lean decision rules

187

AUTOMATED TESTS PASSED

Actual worker and hostile inputs

Detector accuracy remains empirical.

Local demonstration

```
make setup
```

```
make run
```

```
make test
```

Installed local model required.

github.com/ljaniec/Mathguard